

# ASML Project - Exercise 2 - Titanic analysis - José Lise

## - DSTI S19

The goal is to carry out the Titanic data classification analysis. We will use the Titanic dataset available on Kaggle web site. In detail, this is a binary classification problem. The model must be able to predict survival or not with a good accuracy on the test sample.

The data has been splitted into two groups:

- training set (train.csv)
- test set (test.csv)

The *training* set will be used to build the machine learning models. For the training set, is provided the outcome (also known as the “ground truth”) for each passenger.

The *test* set will be used to see how well the models perform on unseen data. For the test set, the ground truth for each passenger is not provided. It is the models’ job to predict these outcomes. For each passenger in the test set, we will use the trained model to predict whether or not they survived the sinking of the Titanic. And as we don’t have the outcome for the test set, we will submit our prediction to the kaggle web site to get our score.

### Loading the data

```
setwd("D:/OneDrive - Data ScienceTech Institute/DSTI/AdvanceStatisticsMachineLearning/Project")
train <- read.csv("titanic/train.csv", stringsAsFactors=FALSE, header=TRUE, sep=',')
test <- read.csv("titanic/test.csv", stringsAsFactors=FALSE, header=TRUE, sep=',')
```

### check the train data frame

```
str(train)

## 'data.frame':    891 obs. of  12 variables:
##  $ PassengerId: int  1 2 3 4 5 6 7 8 9 10 ...
##  $ Survived   : int  0 1 1 1 0 0 0 0 1 1 ...
##  $ Pclass     : int  3 1 3 1 3 3 1 3 3 2 ...
##  $ Name       : chr  "Braund, Mr. Owen Harris" "Cumings, Mrs. John Bradley (Florence Briggs Thayer)"
##  $ Sex        : chr  "male" "female" "female" "female" ...
##  $ Age        : num  22 38 26 35 35 NA 54 2 27 14 ...
##  $ SibSp       : int  1 1 0 1 0 0 0 3 0 1 ...
##  $ Parch       : int  0 0 0 0 0 0 0 1 2 0 ...
##  $ Ticket      : chr  "A/5 21171" "PC 17599" "STON/O2. 3101282" "113803" ...
##  $ Fare        : num  7.25 71.28 7.92 53.1 8.05 ...
##  $ Cabin       : chr  "" "C85" "" "C123" ...
##  $ Embarked    : chr  "S" "C" "S" "S" ...
```

There are 891 observations of 12 variables. 5 variables are integers, 5 are characters and 2 are numeric.

Summary train

```
summary(train)

##  PassengerId      Survived      Pclass         Name
##  Min.         : 1.0      Min.      :0.0000      Min.      :1.000      Length:891
```

```
## 1st Qu.:223.5 1st Qu.:0.0000 1st Qu.:2.000 Class :character
## Median :446.0 Median :0.0000 Median :3.000 Mode :character
## Mean :446.0 Mean :0.3838 Mean :2.309
## 3rd Qu.:668.5 3rd Qu.:1.0000 3rd Qu.:3.000
## Max. :891.0 Max. :1.0000 Max. :3.000
##
## Sex Age SibSp Parch
## Length:891 Min. : 0.42 Min. :0.000 Min. :0.0000
## Class :character 1st Qu.:20.12 1st Qu.:0.000 1st Qu.:0.0000
## Mode :character Median :28.00 Median :0.000 Median :0.0000
## Mean :29.70 Mean :0.523 Mean :0.3816
## 3rd Qu.:38.00 3rd Qu.:1.000 3rd Qu.:0.0000
## Max. :80.00 Max. :8.000 Max. :6.0000
## NA's :177
## Ticket Fare Cabin Embarked
## Length:891 Min. : 0.00 Length:891 Length:891
## Class :character 1st Qu.: 7.91 Class :character Class :character
## Mode :character Median : 14.45 Mode :character Mode :character
## Mean : 32.20
## 3rd Qu.: 31.00
## Max. :512.33
##
```

The summary above already shows that there are 177 missing rows for the age variable.

## Variables description

Here are the short description of the variables in the dataset:

- PassengerId: Identification number for passengers
- Survived: Indicates if the passenger survived: 0=NO, 1=YES
- Pclass: Ticket Class: 1=1st, 2=2nd, 3=3rd
- Sex: Female, Male
- Age: Age in years
- SibSp: # of sibling/Spouses aboard the Titanic
- Parch: # of Parents/Children abroad the Titanic
- Ticket: Ticket number
- fare: Passenger fare
- cabin: Cabin number
- embarked: Port of Embarkation: C=Cherburg, Q=Queenstown, S=Southampton

Here are some additionnal information for the variables: *pclass*: A proxy for socio-economic status (SES)

- 1st = Upper
- 2nd = Middle
- 3rd = Lower

*age*: Age is fractional if less than 1. If the age is estimated, is it in the form of xx.5

*sibsp*: The dataset defines family relations in this way:

- Sibling = brother, sister, stepbrother, stepsister

- Spouse = husband, wife (mistresses and fiancés were ignored)

*parch*: The dataset defines family relations in this way:

- Parent = mother, father
- Child = daughter, son, stepdaughter, stepson  
Some children travelled only with a nanny, therefore parch=0 for them.

check the test data frame

```
str(test)
```

```
## 'data.frame':  418 obs. of  11 variables:
## $ PassengerId: int  892 893 894 895 896 897 898 899 900 901 ...
## $ Pclass     : int   3  3  2  3  3  3  3  2  3  3 ...
## $ Name       : chr   "Kelly, Mr. James" "Wilkes, Mrs. James (Ellen Needs)" "Myles, Mr. Thomas Francis" ...
## $ Sex        : chr   "male" "female" "male" "male" ...
## $ Age        : num   34.5  47  62  27  22  14  30  26  18  21 ...
## $ SibSp      : int    0  1  0  0  1  0  0  1  0  2 ...
## $ Parch      : int    0  0  0  0  1  0  0  1  0  0 ...
## $ Ticket     : chr   "330911" "363272" "240276" "315154" ...
## $ Fare       : num    7.83  7  9.69  8.66  12.29 ...
## $ Cabin      : chr    "" "" "" "" ...
## $ Embarked   : chr    "Q" "S" "Q" "S" ...
```

Test data set contains 418 observation of 11 variables. As expected, the survived variable is missing from this data set.

Test summary

```
summary(test)
```

```
##   PassengerId      Pclass      Name      Sex
##   Min.   : 892.0   Min.   :1.000   Length:418   Length:418
##   1st Qu.: 996.2   1st Qu.:1.000   Class :character   Class :character
##   Median :1100.5   Median :3.000   Mode  :character   Mode  :character
##   Mean   :1100.5   Mean    :2.266
##   3rd Qu.:1204.8   3rd Qu.:3.000
##   Max.   :1309.0   Max.    :3.000
##
##      Age      SibSp      Parch      Ticket
##   Min.   : 0.17   Min.   :0.0000   Min.   :0.0000   Length:418
##   1st Qu.:21.00   1st Qu.:0.0000   1st Qu.:0.0000   Class :character
##   Median :27.00   Median :0.0000   Median :0.0000   Mode  :character
##   Mean   :30.27   Mean    :0.4474   Mean    :0.3923
##   3rd Qu.:39.00   3rd Qu.:1.0000   3rd Qu.:0.0000
##   Max.   :76.00   Max.    :8.0000   Max.    :9.0000
##   NA's    :86
##      Fare      Cabin      Embarked
##   Min.   : 0.000   Length:418   Length:418
##   1st Qu.: 7.896   Class :character   Class :character
##   Median :14.454   Mode  :character   Mode  :character
##   Mean   :35.627
##   3rd Qu.:31.500
##   Max.   :512.329
##   NA's    :1
```

keep raw train, and test data sets for future use during the modeling part. However we will transform the variables Pclass, Sex and Embarked to factors.

```
train_raw <- train
test_raw <- test

train_raw$Pclass <- factor(train_raw$Pclass)
train_raw$Sex <- factor(train_raw$Sex)
train_raw$Embarked <- factor(train_raw$Embarked, exclude="")

test_raw$Pclass <- factor(test_raw$Pclass)
test_raw$Sex <- factor(test_raw$Sex)
test_raw$Embarked <- factor(test_raw$Embarked, exclude="")

test_raw$Survived <- 0
all_raw <- rbind(train_raw, test_raw)
```

Merge train and test data set for exploratory analysis

```
# Create a Survived column for the test dataset and fill it with 0
test$Survived <- 0
all <- rbind(train, test)
```

## Handling Missing Data

```
summary(all)
```

```
## PassengerId      Survived  Pclass      Name
## Min.   :  1    Min.   :0.0000  Min.   :1.000  Length:1309
## 1st Qu.: 328    1st Qu.:0.0000  1st Qu.:2.000  Class :character
## Median : 655    Median :0.0000  Median :3.000  Mode  :character
## Mean   : 655    Mean   :0.2613  Mean   :2.295
## 3rd Qu.: 982    3rd Qu.:1.0000  3rd Qu.:3.000
## Max.   :1309    Max.   :1.0000  Max.   :3.000
##
##      Sex          Age          SibSp          Parch
## Length:1309    Min.   : 0.17    Min.   :0.0000    Min.   :0.000
## Class :character 1st Qu.:21.00    1st Qu.:0.0000    1st Qu.:0.000
## Mode  :character Median :28.00    Median :0.0000    Median :0.000
##                  Mean   :29.88    Mean   :0.4989    Mean   :0.385
##                  3rd Qu.:39.00    3rd Qu.:1.0000    3rd Qu.:0.000
##                  Max.   :80.00    Max.   :8.0000    Max.   :9.000
##                  NA's   :263
##      Ticket      Fare          Cabin
## Length:1309    Min.   : 0.000  Length:1309
## Class :character 1st Qu.: 7.896  Class :character
## Mode  :character Median :14.454  Mode  :character
##                  Mean   :33.295
##                  3rd Qu.:31.275
##                  Max.   :512.329
##                  NA's   :1
##      Embarked
## Length:1309
## Class :character
## Mode  :character
```

```
##
##
##
##
```

In the cell below, we transform the Sex and Embarked variables to factors.

```
all$Sex <- factor(all$Sex)
all$Embarked <- factor(all$Embarked, exclude="")

summary(all)
```

```
## PassengerId      Survived      Pclass      Name
## Min.   : 1      Min.   :0.0000   Min.   :1.000   Length:1309
## 1st Qu.: 328    1st Qu.:0.0000   1st Qu.:2.000   Class :character
## Median : 655    Median :0.0000   Median :3.000   Mode  :character
## Mean   : 655    Mean   :0.2613   Mean   :2.295
## 3rd Qu.: 982    3rd Qu.:1.0000   3rd Qu.:3.000
## Max.   :1309    Max.   :1.0000   Max.   :3.000
##
## Sex              Age              SibSp              Parch
## female:466      Min.   : 0.17      Min.   :0.0000   Min.   :0.000
## male :843       1st Qu.:21.00    1st Qu.:0.0000   1st Qu.:0.000
##                Median :28.00    Median :0.0000   Median :0.000
##                Mean   :29.88    Mean   :0.4989   Mean   :0.385
##                3rd Qu.:39.00    3rd Qu.:1.0000   3rd Qu.:0.000
##                Max.   :80.00    Max.   :8.0000   Max.   :9.000
##                NA's   :263
## Ticket          Fare              Cabin              Embarked
## Length:1309      Min.   : 0.000   Length:1309      C   :270
## Class :character 1st Qu.: 7.896   Class :character  Q   :123
## Mode  :character Median :14.454   Mode  :character  S   :914
##                Mean   :33.295                      NA's: 2
##                3rd Qu.:31.275
##                Max.   :512.329
##                NA's   :1
```

Summary of the missing data

```
sapply(all, function(attribute) {sum(is.na(attribute))==TRUE)/ length(attribute)
;})
```

```
## PassengerId      Survived      Pclass      Name      Sex
## 0.0000000000 0.0000000000 0.0000000000 0.0000000000 0.0000000000
## Age          SibSp          Parch          Ticket          Fare
## 0.2009167303 0.0000000000 0.0000000000 0.0000000000 0.0007639419
## Cabin        Embarked
## 0.0000000000 0.0015278839
```

The output above shows that there are missing values for variables Age, Fare and Embarked. We addressed the missing data for Fare and Embarked in the following way:

- Assign missing Embarked data to the most counted port ('S').
- Replace the missing Fare data by the mean fare.

```
all$Embarked[which(is.na(all$Embarked))] <- 'S'
all$Fare[which(is.na(all$Fare))] <- mean(all$Fare, na.rm=TRUE)
```

```
summary(all)
```

```
## PassengerId      Survived      Pclass      Name
## Min.   :   1      Min.   :0.0000      Min.   :1.000      Length:1309
## 1st Qu.: 328      1st Qu.:0.0000      1st Qu.:2.000      Class :character
## Median : 655      Median :0.0000      Median :3.000      Mode  :character
## Mean   : 655      Mean   :0.2613      Mean   :2.295
## 3rd Qu.: 982      3rd Qu.:1.0000      3rd Qu.:3.000
## Max.   :1309      Max.   :1.0000      Max.   :3.000
##
## Sex              Age              SibSp              Parch
## female:466      Min.   : 0.17      Min.   :0.0000      Min.   :0.000
## male :843      1st Qu.:21.00      1st Qu.:0.0000      1st Qu.:0.000
##              Median :28.00      Median :0.0000      Median :0.000
##              Mean   :29.88      Mean   :0.4989      Mean   :0.385
##              3rd Qu.:39.00      3rd Qu.:1.0000      3rd Qu.:0.000
##              Max.   :80.00      Max.   :8.0000      Max.   :9.000
##              NA's   :263
## Ticket          Fare              Cabin              Embarked
## Length:1309      Min.   : 0.000      Length:1309      C:270
## Class :character  1st Qu.: 7.896      Class :character  Q:123
## Mode  :character  Median :14.454      Mode  :character  S:916
##              Mean   :33.295
##              3rd Qu.:31.275
##              Max.   :512.329
##
```

The Cabin column data is managed as character data. However there are many empty strings. Moreover this variable doesn't provide any relevant information. Therefore we will not use this feature for the modeling part.

```
sum(all$Cabin == "")/nrow(all)
```

```
## [1] 0.7746371
```

There are 77% of empty strings for the Cabin column.

Age Missing Data imputation

1. Check the title frequency

```
table_words = table(unlist(strsplit(all$Name, "\\s+")))
sort(table_words [grep('\\.',names(table_words))], decreasing=TRUE)
```

```
##
## Mr.      Miss.      Mrs.      Master.      Dr.      Rev.      Col.
## 757      260      197      61      8      8      4
## Major.    Mlle.      Ms.      Capt. Countess.      Don.      Dona.
## 2      2      2      1      1      1      1
## Jonkheer. L.      Lady.      Mme.      Sir.
## 1      1      1      1      1
```

2. Find missing age by title

```
library(stringr)
tb_data = cbind(all$Age, str_match(all$Name, "[a-zA-Z]+\\."))
table(tb_data[is.na(tb_data[,1]),2])
```

```
##
##      Dr.  Master.   Miss.    Mr.    Mrs.    Ms.
##      1      8      50     176     27      1
```

3. Compute mean value by titles

```
mean.mr = mean(all$Age[grepl(" Mr\\.", all$Name)],na.rm=TRUE)
mean.mrs = mean(all$Age[grepl(" Mrs\\.", all$Name)],na.rm=TRUE)
mean.dr = mean(all$Age[grepl(" Dr\\.", all$Name)],na.rm=TRUE)
mean.miss = mean(all$Age[grepl(" Miss\\.", all$Name)],na.rm=TRUE)
mean.master = mean(all$Age[grepl(" Master\\.", all$Name)],na.rm=TRUE)
```

4. Apply the mean to the missing data

```
all$Age[grepl(" Mr\\.", all$Name)
      & is.na(all$Age)] = mean.mr
all$Age[grepl(" Mrs\\.", all$Name)
      & is.na(all$Age)] = mean.mrs
all$Age[grepl(" Dr\\.", all$Name)
      & is.na(all$Age)] = mean.dr
all$Age[grepl(" Miss\\.", all$Name)
      & is.na(all$Age)] = mean.miss
all$Age[grepl(" Master\\.", all$Name)
      & is.na(all$Age)] = mean.master
# Special case for Ms. that we manage as Miss.
all$Age[grepl(" Ms\\.", all$Name)
      & is.na(all$Age)] = mean.miss
```

5. Check that there is no remaining missing age values

```
sum(is.na(all$Age) == TRUE) / length(all$Age)
```

```
## [1] 0
```

## Data transformation

Manage class and sex as factors instead of numbers

```
all$Pclass <- factor(all$Pclass)
all$Sex <- factor(all$Sex)
#all$Embarked <- factor(all$Embarked)
```

Add a variable title

```
all$Title <- substring(str_extract(all$Name, '\\, \\w*\\.'), 3)
all$Title <- factor(all$Title)
#all$Title <- gsub('([[:alpha:]]*\\, )([[:alpha:]]*\\.)([[:alpha:]]*)', '\\2', all$Name)

#all$Title <- str_replace(all$Name, '([^\\w*)\\, (\\w*\\.)(\\w*)', REF2)
```


## Exploratory data analysis

Breakdown by gender

```
table(all$Sex)
```

```
##
## female  male
##    466    843
```

```
barplot(table(all$Sex), names= c("Female","Male"), col= c("pink", "lightblue"), main="Breakdown by gender")
```

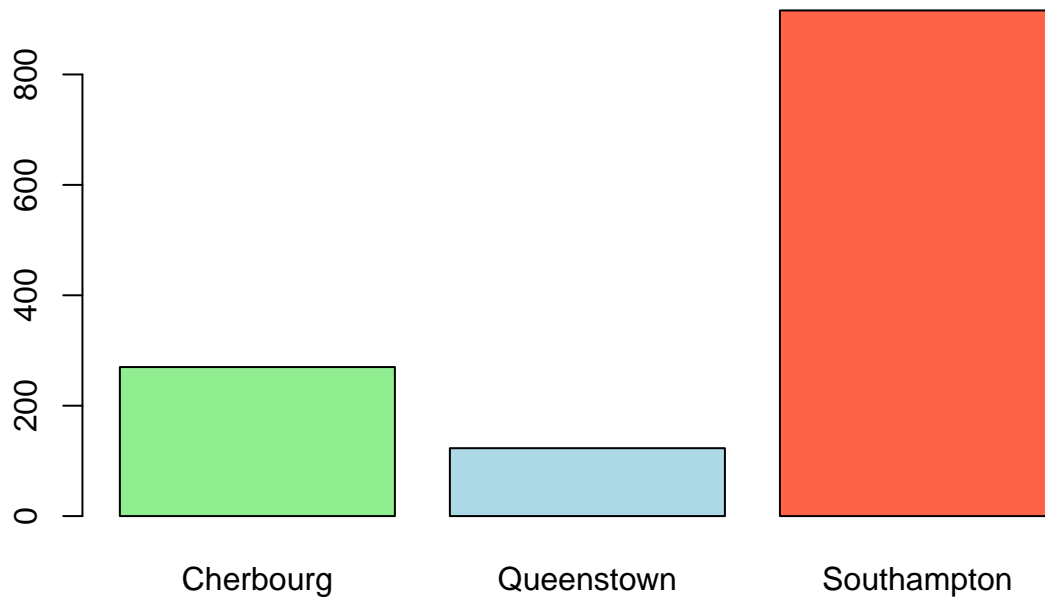


Breakdown by port of Embarkation

```
barplot(table(all$Embarked), col=c("lightgreen","lightblue","tomato"), names= c("Cherbourg", "Queenstown", "Southampton"))
```



## Port of Embarkation



Breakdown by class


```
table(all$Pclass)
```

```
##
```

```
##  1  2  3
```

```
## 323 277 709
```

```
barplot(table(all$Pclass), col=c("lightgreen", "lightblue", "tomato"),  
        names= c("First", "Second", "Third"), main="Breakdown by class")
```




Breakdown by sex for each class

```
table(all$Sex, all$Pclass )
```

```
##  
##           1    2    3  
##  female 144 106 216  
##   male   179 171 493
```

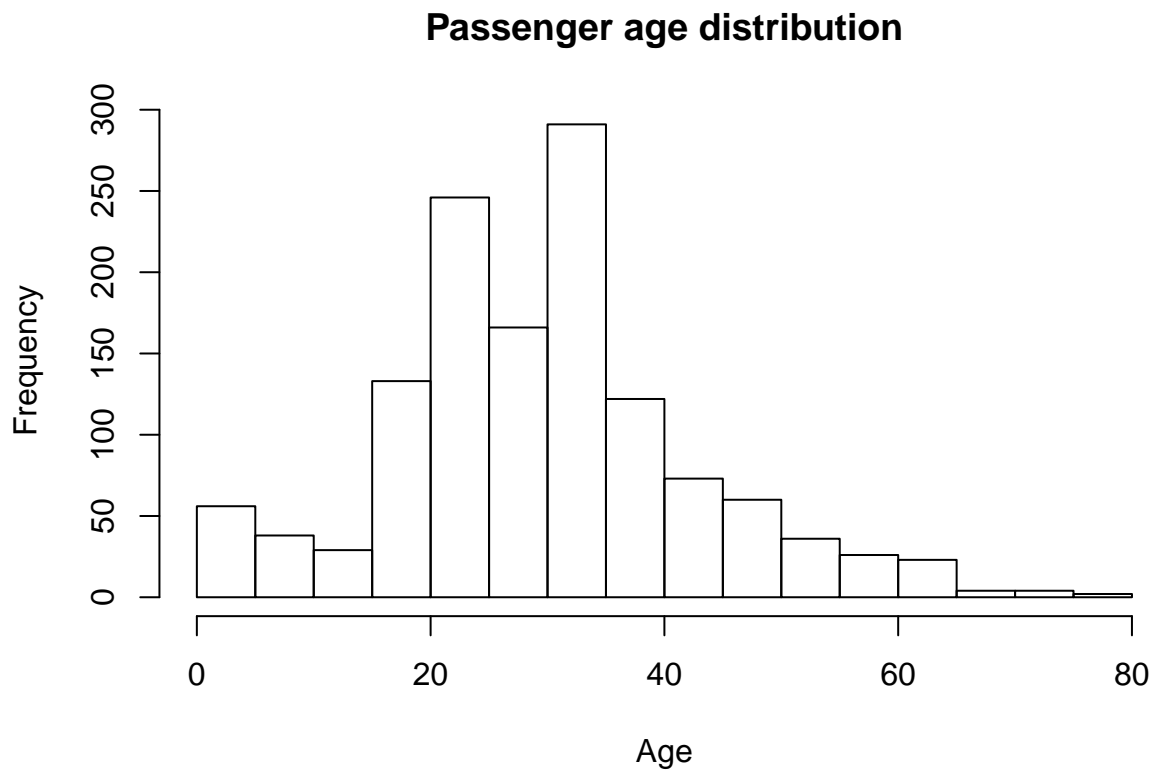
```
countsTable <- table(all$Sex, all$Pclass )  
barplot(countsTable, col=c("pink", "lightblue"), legend=c("Female", "Male"),  
        names=c("First", "Second", "Third"), main="Passengers breakdown by sex for each class")
```

**Passengers breakdown by sex for each class**



Hist distribution by passenger age

```
hist(all$Age, main="Passenger age distribution", xlab="Age")
```




Splitting back the data in train and test sets

```
dt <- 1:nrow(train)
train <- all[dt,]
#mrow_all <- nrow(all)
test <- all[-dt,]
test$Survived <- NULL
```


Specific Exploratory analysis of the training data set

```
barplot(table(train$Survived), col=c("Tomato", "lightgreen"),
        names=c("Perished", "Survived"), legend=c("Perished", "Survived"),
        main="Perished/Survived Breakdown" )
```




Passenger fate by age

```
barplot(table(train$Survived, train$Age), col=c("Tomato", "lightgreen"),  
        legend=c("Perished", "Survived"),  
        main="Passenger fate by age" )
```



Passenger fate by sex

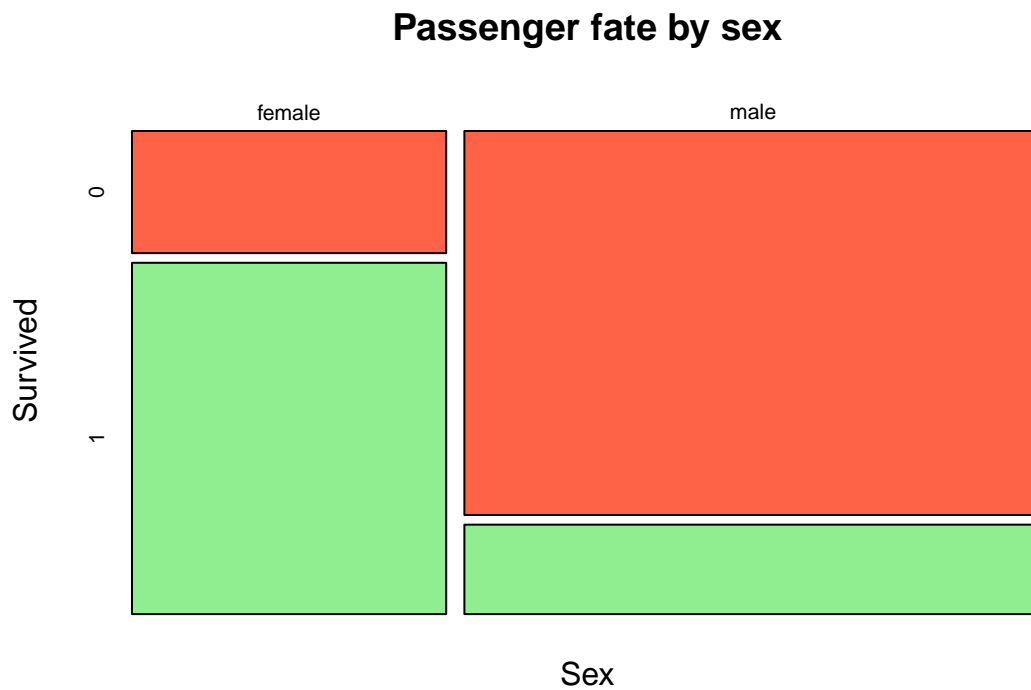
```
barplot(table(train$Survived, train$Sex), col=c("Tomato", "lightgreen"),
        names=c("Female", "Male"), legend=c("Perished", "Survived"),
        main="Passenger fate by sex" )
```



Mosaic plot of the same data

```
mosaicplot( train$Sex~train$Survived, main="Passenger fate by sex",  
            SHADE=FALSE, col=c("Tomato", "lightgreen"),xlab="Sex", ylab="Survived")
```

```
## Warning: In mosaicplot.default(table(mf), main = main, ...) :  
## extra argument 'SHADE' will be disregarded
```



Passenger fate by travelling class

```
table(train$Survived)
```

```
##
##    0    1
## 549 342
```

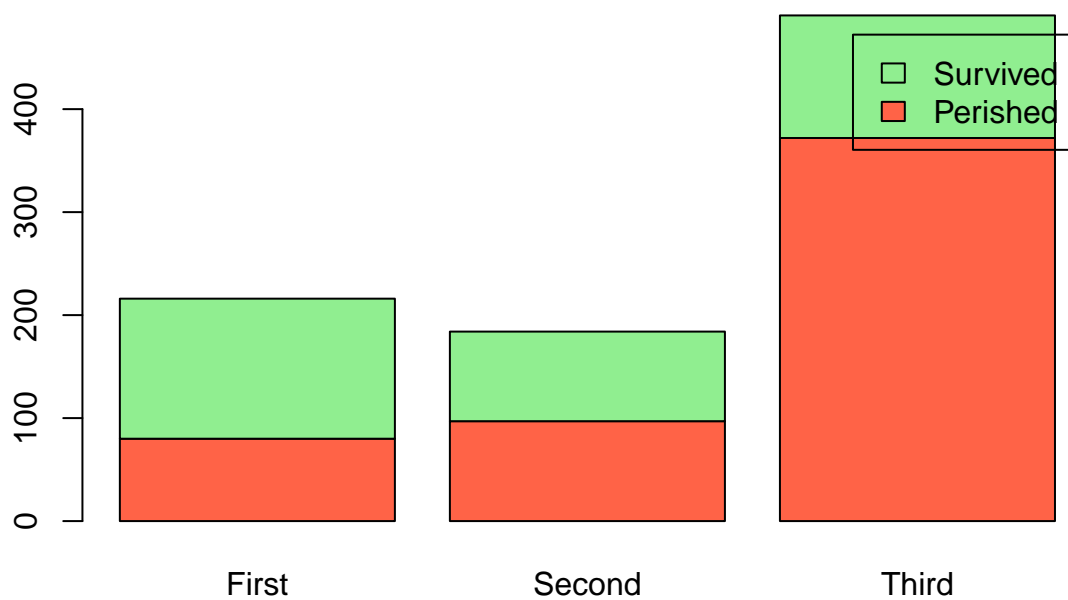
```
table(train$Survived, train$Pclass)
```

```
##
##      1    2    3
##    0  80  97 372
##    1 136  87 119
```

```
barplot( table(train$Survived, train$Pclass), col=c("Tomato", "lightgreen"),
  legend = c("Perished", "Survived"), names= c("First", "Second", "Third"),
  main= "Passenger fate by Class" )
```




## Passenger fate by Class



Corresponding Mosaic Plot

```
mosaicplot(train$Pclass ~ train$Survived, main="Passenger fate by Pclass",  
           shade=FALSE, color=c("tomato","lightgreen"), xlab="Pclass", ylab="Survived")
```



## Predicting passenger survival using Decision Tree

For the modelling part, we will not take into account the following variables:

- PassengerId: This is just an identifier for the passenger and doesn't bring any value.
- Ticket: This is just the ticket number and this doesn't add also any value
- cabin: This is a cabin identification number that is not relevant for this analysis

```
library(rpart)
# Step 1: Build the maximal tree

Tree <- rpart(Survived~Pclass + Sex + Age + SibSp + Parch + Fare + Embarked, data=train,
              method="class", control=rpart.control(minsplit=2,cp=0))

#Tree

Summary r include=FALSE

#summary(Tree)

Error on the maximal tree

pred <- predict(Tree, type="class")
error<- 1/length(train$Survived) * sum(train$Survived != pred )
error


## [1] 0.01795735

printcp
```

```
A <- printcp(Tree)
```

```
##
## Classification tree:
## rpart(formula = Survived ~ Pclass + Sex + Age + SibSp + Parch +
##       Fare + Embarked, data = train, method = "class", control = rpart.control(minsplit = 2,
##       cp = 0))
##
## Variables actually used in tree construction:
## [1] Age      Embarked Fare      Parch    Pclass   Sex      SibSp
##
## Root node error: 342/891 = 0.38384
##
## n= 891
##
##      CP nsplit rel error  xerror    xstd
## 1  0.4444444      0  1.000000 1.00000 0.042446
## 2  0.03070175     1  0.555556 0.55556 0.035750
## 3  0.01461988     5  0.432749 0.48830 0.034061
## 4  0.00730994     6  0.418129 0.49123 0.034140
## 5  0.00584795    10  0.383041 0.50000 0.034372
## 6  0.00487329    12  0.371345 0.47661 0.033744
## 7  0.00438596    22  0.312865 0.46199 0.033336
## 8  0.00292398    29  0.280702 0.46491 0.033419
## 9  0.00194932    50  0.219298 0.45029 0.033001
## 10 0.00167084    56  0.207602 0.48830 0.034061
## 11 0.00146199    65  0.190058 0.52339 0.034970
## 12 0.00125313   122  0.105263 0.53509 0.035260
## 13 0.00116959   137  0.084795 0.54386 0.035472
## 14 0.00109649   142  0.078947 0.54386 0.035472
## 15 0.00097466   153  0.064327 0.56725 0.036021
## 16 0.00073099   159  0.058480 0.56725 0.036021
## 17 0.00058480   169  0.049708 0.56725 0.036021
## 18 0.00000000   174  0.046784 0.56725 0.036021
```

```
plotcp(Tree)
```



Step 2: Pruning

```
mincp <- which(A[,4] == min(A[,4]))
mincp
```

```
## 9
## 9
```

```
#cpthres: 1-SE rule threshold : Error_min + standard_error
cpthres <- A[mincp,4] + A[mincp,5]
cp1se <- min(which(A[,4] <= cpthres))
#cp1se <- which(min(A[cand,4]) == A[,4])
cp1se
```

```
## [1] 6
```

The lower error is 0.44 with a standard error of 0.03. When we apply the 1SE rule, we get  $0.45 + 0.03 = 0.48$  as threshold. Therefore the final tree is the smaller tree (less splits) one with error lower than 0.48. It's the value corresponding to  $cp=0.00487329$  (Id = 6). The code above gives an accurate way to identify the cp corresponding to the 1SE cp rule.

Alternative method

```
cverr=A[,4]
mincverr=which(cverr==min(cverr))
s=A[mincverr,4]+A[mincverr,5]
s=min(s)
B=1*(cverr<=s)
a=min(which(B==1))
```


```

a

## [1] 6
cp=A[a,1]
cp

## [1] 0.004873294
#Treep <- prune(Tree, cp=A[5,1])
Treep <- prune(Tree, cp=A[cp1se,1])
plot(Treep)
text(Treep)

```



Display a more fancy plot

Install package and load library

```

#install.packages('rattle')
#install.packages('rpart.plot')
#install.packages('RColorBrewer')
#library(rattle)

library(rattle)

```

```

## Warning: package 'rattle' was built under R version 3.6.2

## Rattle: A free graphical interface for data science with R.
## Version 5.3.0 Copyright (c) 2006-2018 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.

```

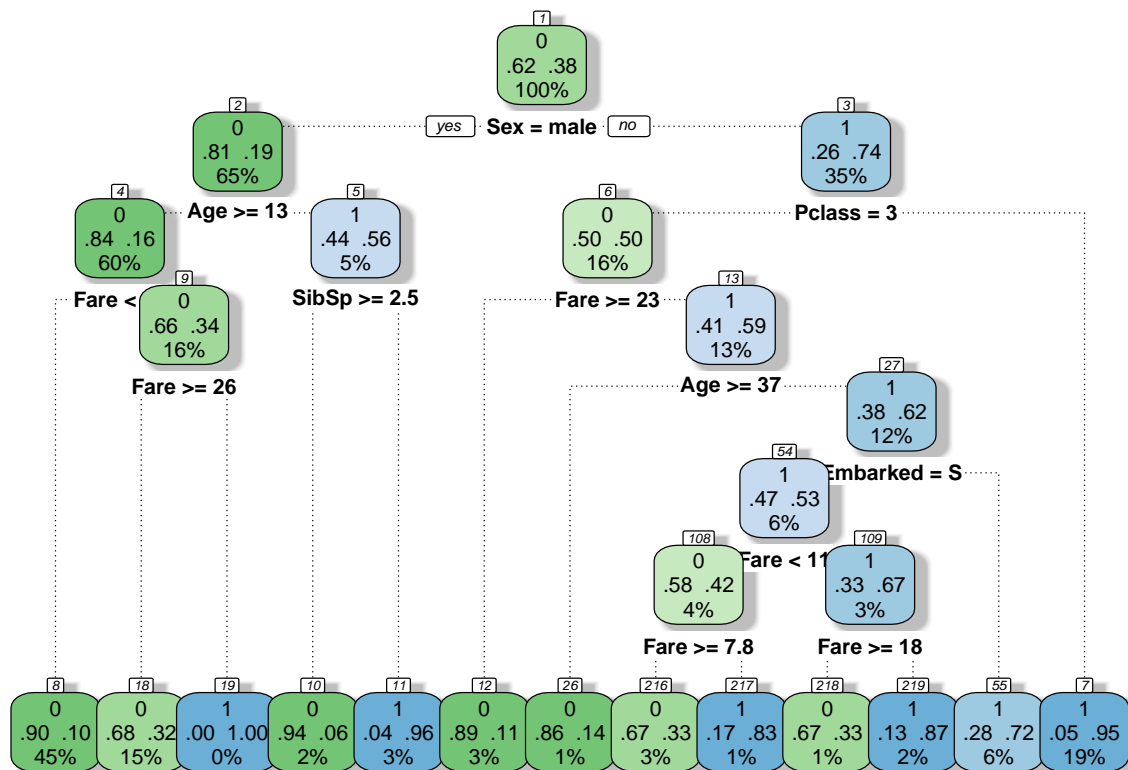
```
library(rpart.plot)
```

```
## Warning: package 'rpart.plot' was built under R version 3.6.2
```

```
library(RColorBrewer)
```

Tree plot

```
fancyRpartPlot(Treeep, tweak=1.4)
```



Rattle 2020-Mar-28 15:06:57 jose

We see from the tree, that the most important variables are respectively: Sex, Age, Pclass and SibSp.

Prediction on the test set

```
pred_dt <- predict(Treeep, newdata=test, type="class")
submit_dt <- data.frame(PassengerId = test$PassengerId, Survived = pred_dt)
```

Write submission

```
write.csv(submit_dt, file = "submit_dt_02.csv", row.names = FALSE)
```

After submission, Kaggle score is 0.79425.

## Random Forest

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.6.1
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
##
## The following object is masked from 'package:rattle':
##
##      importance
# Set seed for reproducibility
set.seed(1234)
Tree_rf <- randomForest(as.factor(Survived) ~ Pclass + Sex + Age + SibSp + Parch +
  Fare + Embarked, data=train, importance=TRUE, proximity=TRUE, ntree=1000)

Tree_rf

##
## Call:
## randomForest(formula = as.factor(Survived) ~ Pclass + Sex + Age + SibSp + Parch + Fare + Embarked,
##               data = train, importance = TRUE, proximity = TRUE, ntree = 1000,
##               type = "classification", number = 1000,
##               variables.tried = 2,
##               oob.error.rate = 0.1728,
##               confusion.matrix = matrix(c(503, 46, 108, 234), nrow = 2, ncol = 2,
##               diag = TRUE),
##               class.error = c(0.08378871, 0.31578947))
##
## No. of variables tried at each split: 2
##
##      OOB estimate of error rate: 17.28%
## Confusion matrix:
##      0      1 class.error
## 0 503  46  0.08378871
## 1 108 234  0.31578947

Prediction

pred_rf <- predict(Tree_rf, newdata=test, type="class")
submit_rf <- data.frame(PassengerId = test$PassengerId, Survived = pred_rf)
```

Write submission

```
write.csv(submit_rf, file = "submit_rf2.csv", row.names = FALSE)
```

The score for Random Forrest is 0.77990 and therefore worse than what we get for decision tree.

## Using CART without prior missing data imputation

We assume here that the train\_raw data.frame is the raw data.frame with missing values:

```
summary(train_raw)
```

```
## PassengerId      Survived  Pclass      Name      Sex
## Min.   : 1.0    Min.   :0.0000  1:216  Length:891  female:314
## 1st Qu.:223.5  1st Qu.:0.0000  2:184  Class :character  male :577
## Median :446.0  Median :0.0000  3:491  Mode  :character
## Mean   :446.0  Mean   :0.3838
## 3rd Qu.:668.5  3rd Qu.:1.0000
## Max.   :891.0  Max.   :1.0000
##
##      Age      SibSp      Parch      Ticket
## Min.   : 0.42  Min.   :0.000  Min.   :0.0000  Length:891
## 1st Qu.:20.12  1st Qu.:0.000  1st Qu.:0.0000  Class :character
## Median :28.00  Median :0.000  Median :0.0000  Mode  :character
## Mean   :29.70  Mean   :0.523  Mean   :0.3816
```

```
## 3rd Qu.:38.00 3rd Qu.:1.000 3rd Qu.:0.0000
## Max. :80.00 Max. :8.000 Max. :6.0000
## NA's :177
## Fare Cabin Embarked
## Min. : 0.00 Length:891 C :168
## 1st Qu.: 7.91 Class :character Q : 77
## Median : 14.45 Mode :character S :644
## Mean : 32.20 NA's: 2
## 3rd Qu.: 31.00
## Max. :512.33
##
```

We confirm that we still have missing values for Age and Embarked fields.

1. Build the maximal tree

```
library(rpart)
# Step 1: Build the maximal tree

Tree_na <- rpart(Survived~Pclass + Sex + Age + SibSp + Parch + Fare + Embarked, data=train_raw,
                 method="class", control=rpart.control(minsplit=2,cp=0))

#Tree_na

printcp
A_na <- printcp(Tree_na)


##
## Classification tree:
## rpart(formula = Survived ~ Pclass + Sex + Age + SibSp + Parch +
##       Fare + Embarked, data = train_raw, method = "class", control = rpart.control(minsplit = 2,
##       cp = 0))
##
## Variables actually used in tree construction:
## [1] Age      Embarked Fare      Parch    Pclass    Sex      SibSp
##
## Root node error: 342/891 = 0.38384
##
## n= 891
##
##      CP nsplit rel error xerror      xstd
## 1  0.4444444      0  1.000000 1.00000 0.042446
## 2  0.03070175     1  0.555556 0.55556 0.035750
## 3  0.02339181     3  0.494152 0.50877 0.034599
## 4  0.02046784     4  0.470760 0.49415 0.034217
## 5  0.00877193     5  0.450292 0.49415 0.034217
## 6  0.00657895    10  0.403509 0.49708 0.034295
## 7  0.00584795    17  0.356725 0.50877 0.034599
## 8  0.00438596    18  0.350877 0.50877 0.034599
## 9  0.00389864    31  0.292398 0.43567 0.032571
## 10 0.00350877    34  0.280702 0.42690 0.032306
## 11 0.00292398    39  0.263158 0.43860 0.032658
## 12 0.00194932    68  0.178363 0.42690 0.032306
## 13 0.00146199    71  0.172515 0.42690 0.032306
## 14 0.00132908   107  0.119883 0.49123 0.034140
```



```
## 15 0.00116959    122 0.099415 0.51170 0.034675
## 16 0.00097466    136 0.078947 0.54971 0.035612
## 17 0.00073099    158 0.055556 0.54971 0.035612
## 18 0.00058480    166 0.049708 0.55848 0.035818
## 19 0.00000000    171 0.046784 0.56140 0.035886
```

```
plotcp
```

```
plotcp(Tree_na)
```



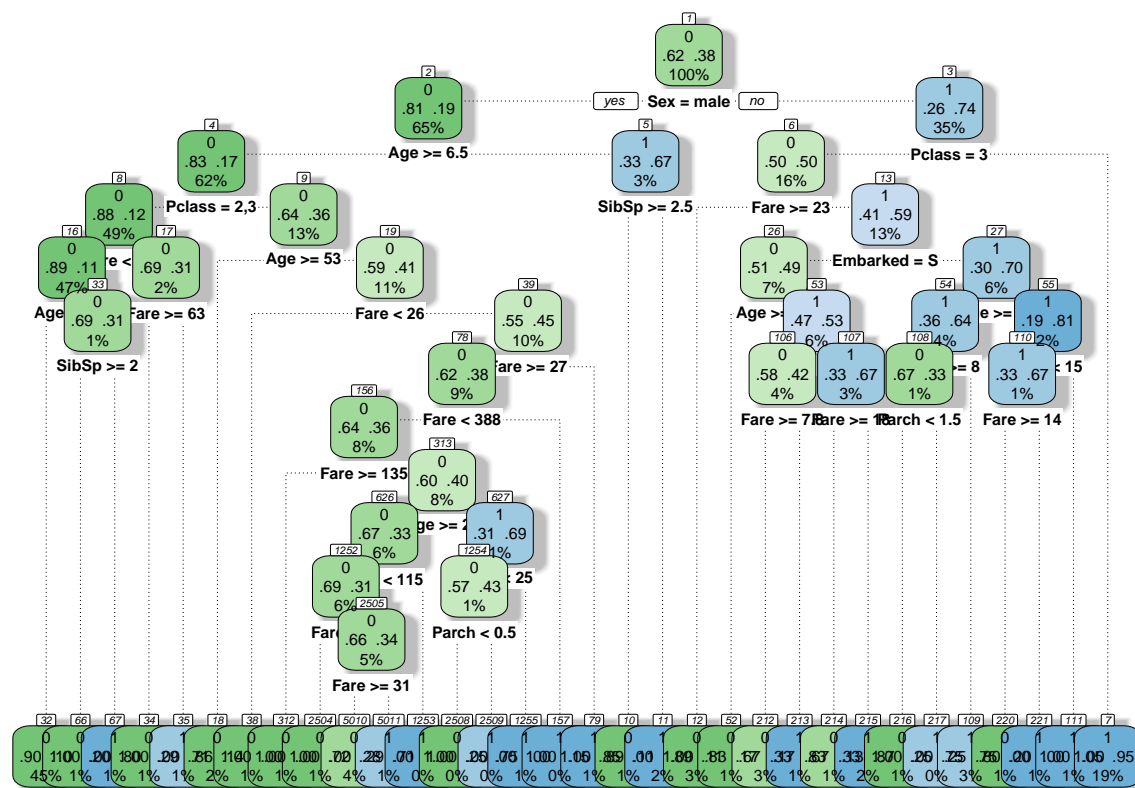
```
mincp <- which(A_na[,4] == min(A_na[,4]))
# as there are several min
mincps <- min(mincp)
#cpthres: 1-SE rule threshold : Error_min + standard_error
cpthres <- A_na[mincps,4] + A_na[mincps,5]
cp1se <- min(which(A_na[,4] <= cpthres))
#cp1se <- which(min(A[cand,4]) == A[,4])
cp1se
```

```
## [1] 9
```

2. Pruning

```
Tree_nap <- prune(Tree_na, cp=A_na[cp1se,1])
plot(Tree_nap)
text(Tree_nap)
```





Rattle 2020-Mar-28 15:07:00 jose

Prediction on the test set

```
pred_dt_na <- predict(Tree_nap, newdata=test_raw, type="class")
submit_dt_na <- data.frame(PassengerId = test_raw$PassengerId, Survived = pred_dt_na)
```

Write submission

```
write.csv(submit_dt_na, file = "submit_dt_na_02.csv", row.names = FALSE)
```

After submission, Kaggle score is 0.77511. Therefore it's worse than model with missing values inputation.

## Using CART with the additional variable title

*# Step 1: Build the maximal tree*

```
Tree_title <- rpart(Survived~Pclass + Sex + Age + SibSp + Parch + Fare + Embarked + Title, data=train,
  method="class", control=rpart.control(minsplit=2,cp=0))
```

*#Tree*

```
A_t1 <- printcp(Tree_title)
```

```
##
```

```
## Classification tree:
```

```
## rpart(formula = Survived ~ Pclass + Sex + Age + SibSp + Parch +
```

```
##   Fare + Embarked + Title, data = train, method = "class",
```

```
##   control = rpart.control(minsplit = 2, cp = 0))
```

```
##
```

```


## Variables actually used in tree construction:
## [1] Age      Embarked Fare      Parch   Pclass   Sex      SibSp    Title
##
## Root node error: 342/891 = 0.38384
##
## n= 891
##
##          CP nsplit rel error  xerror    xstd
## 1  0.46198830      0  1.000000 1.00000 0.042446
## 2  0.05263158      1  0.538012 0.54971 0.035612
## 3  0.01461988      3  0.432749 0.45614 0.033170
## 4  0.00974659      4  0.418129 0.46491 0.033419
## 5  0.00877193      7  0.388889 0.47368 0.033663
## 6  0.00730994      8  0.380117 0.45906 0.033253
## 7  0.00584795     12  0.350877 0.45322 0.033086
## 8  0.00438596     19  0.309942 0.43567 0.032571
## 9  0.00292398     21  0.301170 0.44444 0.032831
## 10 0.00219298     48  0.219298 0.43275 0.032483
## 11 0.00194932     60  0.192982 0.44152 0.032745
## 12 0.00167084     65  0.181287 0.45322 0.033086
## 13 0.00146199     75  0.163743 0.51754 0.034823
## 14 0.00125313    115  0.102339 0.52632 0.035043
## 15 0.00116959    130  0.081871 0.52924 0.035116
## 16 0.00109649    135  0.076023 0.52924 0.035116
## 17 0.00097466    146  0.061404 0.52924 0.035116
## 18 0.00073099    158  0.049708 0.52924 0.035116
## 19 0.00000000    168  0.040936 0.54094 0.035402

```

```

plotcp(Tree_title)

```




```
mincp <- which(A_t1[,4] == min(A_t1[,4]))
# as there are several min
mincps <- min(mincp)
#cpthres: 1-SE rule threshold : Error_min + standard_error
cpthres <- A_t1[mincps,4] + A_t1[mincps,5]
cp1se <- min(which(A_t1[,4] <= cpthres))
#cp1se <- which(min(A[cand,4]) == A[,4])
cp1se
```

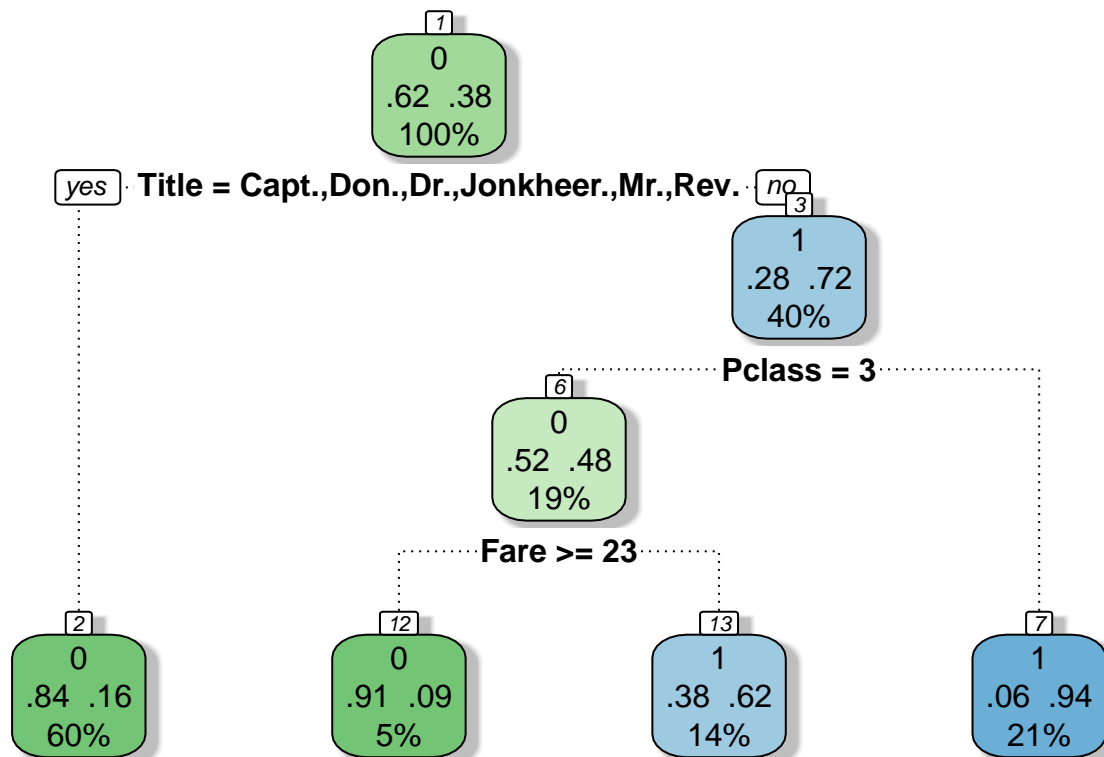
```
## [1] 3
```

```
Pruning
```

```
Tree_tlp <- prune(Tree_title, cp=A_t1[cp1se,1])
plot(Tree_tlp)
text(Tree_tlp)
```



```
fancyRpartPlot(Tree_tlp, tweak=1.0)
```



Rattle 2020-Mar-28 15:07:00 jose

Prediction on the test set

```
# Add an extra factor Dona. for title Dona.
#levels(train$Title) <- c(levels(train$Title), "Dona.")
pred_dt_t1 <- predict(Tree_tlp, newdata=test, type="class")
submit_dt_t1 <- data.frame(PassengerId = test$PassengerId, Survived = pred_dt_t1)
```

Write submission

```
write.csv(submit_dt_t1, file = "submit_dt_t1_01.csv", row.names = FALSE)
```

After submission, Kaggle score is 0.79425

## Conclusion

We tried to predict the fate of passengers in the test set using several CART and Random Forrest models. We got the best score with the CART model with simple inputations for missing data. For this specific problem we were not able to get better results with Random Forrest models.